

# Package ‘EmpiricalCalibration’

January 13, 2020

**Type** Package

**Title** Routines for Performing Empirical Calibration of Observational Study Estimates

**Version** 2.0.1

**Date** 2019-07-19

**Maintainer** Martijn Schuemie <schuemie@ohdsi.org>

**Description** Routines for performing empirical calibration of observational study estimates. By using a set of negative control hypotheses we can estimate the empirical null distribution of a particular observational study setup. This empirical null distribution can be used to compute a calibrated p-value, which reflects the probability of observing an estimated effect size when the null hypothesis is true taking both random and systematic error into account. A similar approach can be used to calibrate confidence intervals, using both negative and positive controls.

**VignetteBuilder** knitr

**Imports** ggplot2 (>= 2.0.0),  
gridExtra,  
methods

**Suggests** knitr,  
rmarkdown

**License** Apache License 2.0

**URL** <https://ohdsi.github.io/EmpiricalCalibration>, <https://github.com/OHDSI/EmpiricalCalibration>

**BugReports** <https://github.com/OHDSI/EmpiricalCalibration/issues>

**RoxygenNote** 7.0.2

## R topics documented:

calibrateConfidenceInterval . . . . .	2
calibrateP . . . . .	3
caseControl . . . . .	4
cohortMethod . . . . .	5
computeTraditionalCi . . . . .	6
computeTraditionalP . . . . .	6
convertNullToErrorModel . . . . .	7
EmpiricalCalibration . . . . .	8

evaluateCiCalibration . . . . .	8
fitMcmcNull . . . . .	9
fitNull . . . . .	10
fitSystematicErrorModel . . . . .	11
grahamReplication . . . . .	12
plotCalibration . . . . .	13
plotCalibrationEffect . . . . .	14
plotCiCalibration . . . . .	15
plotCiCalibrationEffect . . . . .	16
plotCiCoverage . . . . .	17
plotErrorModel . . . . .	18
plotExpectedTypeIError . . . . .	19
plotForest . . . . .	21
plotMcmcTrace . . . . .	22
plotTrueAndObserved . . . . .	22
sccs . . . . .	23
simulateControls . . . . .	24
southworthReplication . . . . .	25

## Index 26

---

calibrateConfidenceInterval  
*Calibrate confidence intervals*

---

### Description

Calibrate confidence intervals

### Usage

```
calibrateConfidenceInterval(logRr, seLogRr, model, ciWidth = 0.95)
```

### Arguments

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = (log(<lower bound 95 percent confidence interval>) - log(<effect estimate>))/qnorm(0.025).
model	An object of type systematicErrorModel as created by the <a href="#">fitSystematicErrorModel</a> function.
ciWidth	The width of the confidence interval. Typically this would be .95, for the 95 percent confidence interval.

### Details

Compute calibrated confidence intervals based on a model of the systematic error.

### Value

A data frame with calibrated confidence intervals and point estimates.

**Examples**

```
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
model <- fitSystematicErrorModel(data$logRr, data$seLogRr, data$trueLogRr)
newData <- simulateControls(n = 15, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
result <- calibrateConfidenceInterval(newData$logRr, newData$seLogRr, model)
result
```

calibrateP

*Calibrate the p-value***Description**

calibrateP computes calibrated p-values using the fitted null distribution

**Usage**

```
calibrateP(null, logRr, seLogRr, ...)

## S3 method for class 'null'
calibrateP(null, logRr, seLogRr, ...)

## S3 method for class 'mcmcNull'
calibrateP(null, logRr, seLogRr, pValueOnly, ...)
```

**Arguments**

null	An object of class null created using the fitNull function or an object of class mcmcNull created using the fitMcmcNull function.
logRr	A numeric vector of one or more effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = (log(<lower bound 95 percent confidence interval>) - log(<effect estimate>))/qnorm(0.025)
...	Any additional parameters (currently none).
pValueOnly	If true, will return only the calibrated P-value itself, not the credible interval.

**Details**

This function computes a calibrated two-sided p-value as described in Schuemie et al (2014).

**Value**

The two-sided calibrated p-value.

**Methods (by class)**

- null: Computes the calibrated P-value using asymptotic assumptions.
- mcmcNull: Computes the calibrated P-value and 95 percent credible interval using Markov Chain Monte Carlo (MCMC).

## References

Schuemie MJ, Ryan PB, Dumouchel W, Suchard MA, Madigan D. Interpreting observational studies: why empirical calibration is needed to correct p-values. *Statistics in Medicine* 33(2):209-18,2014

## Examples

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
null <- fitNull(negatives$logRr, negatives$seLogRr)
positive <- sccs[sccs$groundTruth == 1, ]
calibrateP(null, positive$logRr, positive$seLogRr)
```

---

caseControl

*Odds ratios from a case-control design*

---

## Description

Odds ratios from a case-control design

## Usage

```
data(caseControl)
```

## Format

A data frame with 47 rows and 4 variables:

**drugName** Name of the drug

**groundTruth** Whether the drug is a positive (1) or negative (0) control

**logRr** The log of the incidence rate ratio

**seLogRr** The standard error of the log of the incidence rate ratio

## Details

A dataset containing the odds ratios (and standard errors) produced using a case-control design. The outcome is upper GI bleeding, the drug of interest (`groundTruth = 1`) is sertraline. Also included are 46 negative control drugs, for which we believe there to be no causal relation with upper GI bleeding. We used a database of medical records from general practices in the USA, the General Electric (GE) Centricity database, which contains data on 11.2 million subjects. We restricted on study period (start of 1990 through November 2003), age requirements (18 years or older), available time prior to event (180 days), number of controls per case (6), and risk definition window (30 days following the prescription). Controls were matched on age and sex. Cases of upper GI bleeding were identified on the basis of the occurrence of ICD-9 diagnosis codes in the problem list. These codes pertain to esophageal, gastric, duodenal, peptic, and gastrojejunal ulceration, perforation, and hemorrhage, as well as gastritis and non-specific gastrointestinal hemorrhage. For more information on this set see Schuemie et al (2014).

## References

Schuemie MJ, Ryan PB, Dumouchel W, Suchard MA, Madigan D. Interpreting observational studies: why empirical calibration is needed to correct p-values. *Statistics in Medicine* 33(2):209-18,2014

---

cohortMethod

*Relative risks from a new-user cohort design*


---

## Description

Relative risks from a new-user cohort design

## Usage

```
data(cohortMethod)
```

## Format

A data frame with 31 rows and 4 variables:

**drugName** Name of the drug

**groundTruth** Whether the drug is a positive (1) or negative (0) control

**logRr** The log of the incidence rate ratio

**seLogRr** The standard error of the log of the incidence rate ratio

## Details

A dataset containing the relative risks (and standard errors) produced using a new-user cohort design. The outcome is acute liver injury, the drug of interest (groundTruth = 1) is Isoniazid. Also included are 30 negative control drugs, for which we believe there to be no causal relation with acute liver injury. We used the Thomson MarketScan Medicare Supplemental Beneficiaries database, which contains data on 4.6 million subjects. We selected two groups (cohorts): (1) all subjects exposed to isoniazid and (2) all subjects having the ailment for which isoniazid is indicated, in this case tuberculosis, and having received at least one drug that is not known to cause acute liver injury. We removed all subjects who belonged to both groups and subjects for which less than 180 days of observation time was available prior to their first exposure to the drug in question. Acute liver injury was identified on the basis of the occurrence of ICD-9-based diagnosis codes from inpatient and outpatient medical claims and was defined broadly on the basis of codes associated with hepatic dysfunction, as have been used in prior observational database studies. The time at risk was defined as the length of exposure + 30 days, and we determined whether subjects experienced an acute liver injury during their time at risk. Using propensity score stratification, the cohorts were divided over 20 strata, and an odds ratio over all strata was computed using a Mantel-Haenszel test. The propensity score was estimated using Bayesian logistic regression using all available drug, condition, and procedure covariates occurring in the 180 days prior to first exposure, in addition to age, sex, calendar year of first exposure, Charlson index, number of drugs, number of visit days, and number of procedures. For more information on this set see Schuemie et al (2014).

## References

Schuemie MJ, Ryan PB, Dumouchel W, Suchard MA, Madigan D. Interpreting observational studies: why empirical calibration is needed to correct p-values. *Statistics in Medicine* 33(2):209-18,2014

---

computeTraditionalCi    *Compute the (traditional) confidence interval*

---

### Description

computeTraditionalCi computes the traditional confidence interval based on the log of the relative risk and the standard error of the log of the relative risk.

### Usage

```
computeTraditionalCi(logRr, seLogRr, ciWidth = 0.95)
```

### Arguments

logRr	A numeric vector of one or more effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$
ciWidth	The width of the confidence interval. Typically this would be .95, for the 95 percent confidence interval.

### Value

The point estimate and confidence interval

### Examples

```
data(sccs)
positive <- sccs[sccs$groundTruth == 1, ]
computeTraditionalCi(positive$logRr, positive$seLogRr)
```

---

computeTraditionalP    *Compute the (traditional) p-value*

---

### Description

computeTraditionalP computes the traditional two-sided p-value based on the log of the relative risk and the standard error of the log of the relative risk.

### Usage

```
computeTraditionalP(logRr, seLogRr)
```

### Arguments

logRr	A numeric vector of one or more effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$

**Value**

The two-sided (traditional) p-value.

**Examples**

```
data(sccs)
positive <- sccs[sccs$groundTruth == 1, ]
computeTraditionalP(positive$logRr, positive$sseLogRr)
```

---

```
convertNullToErrorModel
```

*Convert empirical null distribution to systematic error model*

---

**Description**

This function converts an empirical null distribution, fitted using estimates only for negative controls, into a systematic error distribution that can be used to calibrate confidence intervals in addition to p-values.

Whereas the [fitSystematicErrorModel](#) uses positive controls to determine how the error distribution changes with true effect size, this function requires the user to make an assumption. The default assumption, `meanSlope = 1` and `sdSlope = 0`, specify a belief that the error distribution is the same for all true effect sizes. In many cases this assumption is likely to be correct, however, if an estimation method is biased towards the null this assumption will be violated, causing the calibrated confidence intervals to have lower than nominal coverage.

**Usage**

```
convertNullToErrorModel(null, meanSlope = 1, sdSlope = 0)
```

**Arguments**

<code>null</code>	The empirical null distribution fitted using either the <a href="#">fitNull</a> or the <a href="#">fitMcmcNull</a> function.
<code>meanSlope</code>	The slope for the mean of the error distribution. A slope of 1 means the error is the same for different values of the true relative risk.
<code>sdSlope</code>	The slope for the log of the standard deviation of the error distribution. A slope of 0 means the standard deviation is the same for different values of the true relative risk.

**Value**

An object of type `systematicErrorModel`.

**Examples**

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
null <- fitNull(negatives$logRr, negatives$sseLogRr)
model <- convertNullToErrorModel(null)
positive <- sccs[sccs$groundTruth == 1, ]
calibrateConfidenceInterval(positive$logRr, positive$sseLogRr, model)
```

---

EmpiricalCalibration    *EmpiricalCalibration*

---

## Description

EmpiricalCalibration

---

evaluateCiCalibration    *Evaluate confidence interval calibration*

---

## Description

evaluateCiCalibration performs a leave-one-out cross-validation to evaluate the calibration confidence intervals.

## Usage

```
evaluateCiCalibration(
  logRr,
  seLogRr,
  trueLogRr,
  strata = as.factor(trueLogRr),
  crossValidationGroup = 1:length(logRr),
  legacy = FALSE
)
```

## Arguments

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = (log(<lower bound 95 percent confidence interval>) - log(<effect estimate>))/qnorm(0.025).
trueLogRr	The true log relative risk.
strata	Variable used to stratify the plot. Set strata = NULL for no stratification.
crossValidationGroup	What should be the unit for the cross-validation? By default the unit is a single control, but a different grouping can be provided, for example linking a negative control to synthetic positive controls derived from that negative control.
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.

## Details

The empirical calibration is performed using a leave-one-out design: The confidence interval of an effect is computed by fitting a null using all other controls.



**Value**

A data frame specifying the coverage per strata (usually true effect size) for a wide range of widths of the confidence interval. The result also includes the fraction of estimates that was below and above the confidence interval.

**Examples**

```
## Not run:
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
eval <- evaluateCiCalibration(data$logRr, data$seLogRr, data$trueLogRr)

## End(Not run)
```

fitMcmcNull

*Fit the null distribution using MCMC***Description**

fitNull fits the null distribution to a set of negative controls using Markov Chain Monte Carlo (MCMC).

**Usage**

```
fitMcmcNull(logRr, seLogRr, iter = 10000)
```

**Arguments**

logRr	A numeric vector of effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / qnorm(0.025)$
iter	Number of iterations of the MCMC.

**Details**

This is an experimental function for computing the 95 percent credible interval of a calibrated p-value using Markov-Chain Monte Carlo (MCMC).

**Value**

An object of type `mcmcNull` containing the mean and standard deviation (both on the log scale) of the null distribution, as well as the MCMC trace.

**Examples**

```
## Not run:
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
null <- fitMcmcNull(negatives$logRr, negatives$seLogRr)
null
plotMcmcTrace(null)
positive <- sccs[sccs$groundTruth == 1, ]
```

```
calibrateP(null, positive$logRr, positive$seLogRr)

## End(Not run)
```

---

fitNull

*Fit the null distribution*


---

## Description

fitNull fits the null distribution to a set of negative controls

## Usage

```
fitNull(logRr, seLogRr)
```

## Arguments

logRr	A numeric vector of effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$

## Details

This function fits a Gaussian function to the negative control estimates as described in Schuemie et al (2014).

## Value

An object containing the parameters of the null distribution.

## References

Schuemie MJ, Ryan PB, Dumouchel W, Suchard MA, Madigan D. Interpreting observational studies: why empirical calibration is needed to correct p-values. *Statistics in Medicine* 33(2):209-18,2014

## Examples

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
null <- fitNull(negatives$logRr, negatives$seLogRr)
null
```

---

fitSystematicErrorModel

*Fit a systematic error model*


---

## Description

Fit a systematic error model

## Usage

```
fitSystematicErrorModel(
  logRr,
  seLogRr,
  trueLogRr,
  estimateCovarianceMatrix = FALSE,
  legacy = FALSE
)
```

## Arguments

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = (log(<lower bound 95 percent confidence interval>) - log(<effect estimate>))/qnorm(0.025).
trueLogRr	A vector of the true effect sizes.
estimateCovarianceMatrix	Should a covariance matrix be computed? If so, confidence intervals for the model parameters will be available.
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.

## Details

Fit a model of the systematic error as a function of true effect size. This model is an extension of the method for fitting the null distribution. The mean and log(standard deviations) of the error distributions are assumed to be linear with respect to the true effect size, and each component is therefore represented by an intercept and a slope.

## Value

An object of type systematicErrorModel.

## Examples

```
controls <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
model <- fitSystematicErrorModel(controls$logRr, controls$seLogRr, controls$trueLogRr)
model
```

---

grahamReplication	<i>Relative risks from an adjusted new-user cohort design</i>
-------------------	---

---

## Description

Relative risks from an adjusted new-user cohort design

## Usage

```
data(grahamReplication)
```

## Format

A data frame with 126 rows and 4 variables:

**outcomeName** Name of the outcome

**trueLogRr** The log of the true effect size. Only provided for negative and positive controls, is NA for the outcome of interest (GI bleeding).

**logRr** The log of the incidence rate ratio

**seLogRr** The standard error of the log of the incidence rate ratio

## Details

A dataset containing the incidence rate ratios (and standard errors) produced using a new-user cohort design that compares new-users of dabigatran to new-users of warfarin for the outcome of GI hemorrhage. The dataset includes estimates both for the outcome of interest as well as negative and positive control outcomes. Subject are required to have 183 days of continuous observation prior to initiating treatment, be at least 65 years old at index date, and are required to have no prior exposure to warfarin or dabigatran (or any other novel anticoagulant). Furthermore, subjects are required to use the treatment for the indication of atrial fibrillation or atrial flutter, which is enforced by requiring a prior diagnosis of atrial fibrillation or flutter, and no prior diagnosis of other indications. Propensity scores are generated by fitting a model for predicting treatment assignment based on baseline patient characteristics, and are used to perform one-on-one matching. Hazard ratios are estimated through a Cox regression on the matched population. Time-at-risk is defined as starting on the day after initiating treatment and stopping when treatment is stopped, when the outcome occurs, or observation time ends, whichever comes first. The original study (Graham et al 2016) uses the Medicare database. For our replication, we use the Truven Medicare Supplementary Beneficiaries database. We analyze 15,796 dabigatran-exposed and 15,796 warfarin-exposed subjects. For more information on this set see Schuemie et al (2017).

## References

Schuemie MJ, Hripcsak G, Ryan PB, Madigan D, Suchard MA. Empirical confidence interval calibration for population-level effect estimation studies in observational healthcare data. *Proc Natl Acad Sci U S A*. 2018 Mar 13;115(11):2571-2577

Graham DJ, Reichman ME, Wernecke M, Hsueh YH, Izem R, Southworth MR, Wei Y, Liao J, Goulding MR, Mott K, Chillarige Y, MaCurdy TE, Worrall C, Kelman JA. Stroke, Bleeding, and Mortality Risks in Elderly Medicare Beneficiaries Treated With Dabigatran or Rivaroxaban for Nonvalvular Atrial Fibrillation. *JAMA Intern Med* 176(11):1662-1671, 2016

---

plotCalibration	<i>Create a calibration plot</i>
-----------------	----------------------------------

---

## Description

plotCalibration creates a plot showing the calibration of our calibration procedure

## Usage

```
plotCalibration(
  logRr,
  seLogRr,
  useMcmc = FALSE,
  legendPosition = "right",
  title,
  fileName = NULL
)
```

## Arguments

logRr	A numeric vector of effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = (log(<lower bound 95 percent confidence interval>) - log(<effect estimate>))/qnorm(0.025)
useMcmc	Use MCMC to estimate the calibrated P-value?
legendPosition	Where should the legend be positioned? ("none", "left", "right", "bottom", "top")
title	Optional: the main title for the plot
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

## Details

Creates a calibration plot showing the number of effects with  $p < \alpha$  for every level of  $\alpha$ . The empirical calibration is performed using a leave-one-out design: The p-value of an effect is computed by fitting a null using all other negative controls. Ideally, the calibration line should approximate the diagonal. The plot shows both theoretical (traditional) and empirically calibrated p-values.

## Value

A Ggplot object. Use the ggsave function to save to file.

## Examples

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
plotCalibration(negatives$logRr, negatives$seLogRr)
```

---

plotCalibrationEffect *Plot the effect of the calibration*

---

## Description

plotCalibrationEffect creates a plot showing the effect of the calibration.

## Usage

```
plotCalibrationEffect(
  logRrNegatives,
  seLogRrNegatives,
  logRrPositives = NULL,
  seLogRrPositives = NULL,
  null = NULL,
  alpha = 0.05,
  xLabel = "Relative risk",
  title,
  showCis = FALSE,
  fileName = NULL
)
```

## Arguments

logRrNegatives	A numeric vector of effect estimates of the negative controls on the log scale.
seLogRrNegatives	The standard error of the log of the effect estimates of the negative controls.
logRrPositives	Optional: A numeric vector of effect estimates of the positive controls on the log scale.
seLogRrPositives	Optional: The standard error of the log of the effect estimates of the positive controls.
null	An object representing the fitted null distribution as created by the fitNull or fitMcmcNull functions. If not provided, a null will be fitted before plotting.
alpha	The alpha for the hypothesis test.
xLabel	The label on the x-axis: the name of the effect estimate.
title	Optional: the main title for the plot
showCis	Show 95 percent credible intervals for the calibrated $p = \alpha$ boundary.
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

## Details

Creates a plot with the effect estimate on the x-axis and the standard error on the y-axis. Negative controls are shown as blue dots, positive controls as yellow diamonds. The area below the dashed line indicated estimates with  $p < 0.05$ . The orange area indicates estimates with calibrated  $p < 0.05$ .

## Value

A Ggplot object. Use the ggsave function to save to file.

**Examples**

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
positive <- sccs[sccs$groundTruth == 1, ]
plotCalibrationEffect(negatives$logRr, negatives$seLogRr, positive$logRr, positive$seLogRr)
```

---

plotCiCalibration	Create a confidence interval calibration plot
-------------------	---

---

**Description**

plotCalibration creates a plot showing the calibration of our confidence interval calibration procedure

**Usage**

```
plotCiCalibration(
  logRr,
  seLogRr,
  trueLogRr,
  strata = as.factor(trueLogRr),
  crossValidationGroup = 1:length(logRr),
  legacy = FALSE,
  evaluation,
  legendPosition = "top",
  title,
  fileName = NULL
)
```

**Arguments**

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$ .
trueLogRr	The true log relative risk.
strata	Variable used to stratify the plot. Set strata = NULL for no stratification.
crossValidationGroup	What should be the unit for the cross-validation? By default the unit is a single control, but a different grouping can be provided, for example linking a negative control to synthetic positive controls derived from that negative control.
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.
evaluation	A data frame as generated by the <a href="#">evaluateCiCalibration</a> function. If provided, the logRr, seLogRr, trueLogRr, strata, and legacy arguments will be ignored.
legendPosition	Where should the legend be positioned? ("none", "left", "right", "bottom", "top").

title	Optional: the main title for the plot
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

### Details

Creates a calibration plot showing the fraction of effects within the confidence interval. The empirical calibration is performed using a leave-one-out design: The confidence interval of an effect is computed by fitting a null using all other controls. Ideally, the calibration line should approximate the diagonal. The plot shows the coverage for both theoretical (traditional) and empirically calibrated confidence intervals.

### Value

A Ggplot object. Use the ggsave function to save to file.

### Examples

```
## Not run:
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotCiCalibration(data$logRr, data$sseLogRr, data$trueLogRr)

## End(Not run)
```

---

```
plotCiCalibrationEffect
```

*Plot true and observed values*

---

### Description

Creates a plot with the effect estimate on the x-axis and the standard error on the y-axis. The plot is trellised by true effect size. Negative and positive controls are shown as blue dots. The area below the dashed line indicated estimates that are statistically significant different from the true effect size ( $p < 0.05$ ). The orange area indicates estimates with calibrated  $p < 0.05$ .

### Usage

```
plotCiCalibrationEffect(
  logRr,
  seLogRr,
  trueLogRr,
  legacy = FALSE,
  model = NULL,
  xLabel = "Relative risk",
  title,
  fileName = NULL
)
```



**Arguments**

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$ .
trueLogRr	A vector of the true effect sizes.
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.
model	The fitted systematic error model. If not provided, it will be fitted on the provided data.
xLabel	The label on the x-axis: the name of the effect estimate.
title	Optional: the main title for the plot
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

**Value**

A Ggplot object. Use the ggsave function to save to file.

**Examples**

```
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotCiCalibrationEffect(data$logRr, data$seLogRr, data$trueLogRr)
```

---

plotCiCoverage	<i>Create a confidence interval coverage plot</i>
----------------	---

---

**Description**

plotCiCoverage creates a plot showing the coverage before and after confidence interval calibration at various widths of the confidence interval.

**Usage**

```
plotCiCoverage(
  logRr,
  seLogRr,
  trueLogRr,
  strata = as.factor(trueLogRr),
  crossValidationGroup = 1:length(logRr),
  legacy = FALSE,
  evaluation,
  legendPosition = "top",
  title,
  fileName = NULL
)
```

**Arguments**

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$ .
trueLogRr	The true log relative risk.
strata	Variable used to stratify the plot. Set strata = NULL for no stratification.
crossValidationGroup	What should be the unit for the cross-validation? By default the unit is a single control, but a different grouping can be provided, for example linking a negative control to synthetic positive controls derived from that negative control.
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.
evaluation	A data frame as generated by the <a href="#">evaluateCiCalibration</a> function. If provided, the logRr, seLogRr, trueLogRr, strata, and legacy arguments will be ignored.
legendPosition	Where should the legend be positioned? ("none", "left", "right", "bottom", "top").
title	Optional: the main title for the plot
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

**Details**

Creates a plot showing the fraction of effects above, within, and below the confidence interval. The empirical calibration is performed using a leave-one-out design: The confidence interval of an effect is computed by fitting a null using all other controls. The plot shows the coverage for both theoretical (traditional) and empirically calibrated confidence intervals.

**Value**

A Ggplot object. Use the ggsave function to save to file.

**Examples**

```
## Not run:
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotCiCoverage(data$logRr, data$seLogRr, data$trueLogRr)

## End(Not run)
```

---

plotErrorModel

*Plot the systematic error model*


---

**Description**

plotErrorModel creates a plot showing the systematic error model.

**Usage**

```
plotErrorModel(
  logRr,
  seLogRr,
  trueLogRr,
  title,
  legacy = FALSE,
  fileName = NULL
)
```

**Arguments**

logRr	A numeric vector of effect estimates on the log scale.
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / qnorm(0.025)$ .
trueLogRr	The true log relative risk.
title	Optional: the main title for the plot
legacy	If true, a legacy error model will be fitted, meaning standard deviation is linear on the log scale. If false, standard deviation is assumed to be simply linear.
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

**Details**

Creates a plot with the true effect size on the x-axis, and the mean plus and minus the standard deviation shown on the y-axis. Also shown are simple error models fitted at each true relative risk in the input.

**Value**

A Ggplot object. Use the ggsave function to save to file.

**Examples**

```
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotErrorModel(data$logRr, data$seLogRr, data$trueLogRr)
```

---

plotExpectedType1Error

*Plot the expected type 1 error as a function of standard error*

---

**Description**

plotExpectedType1Error creates a plot showing the expected type 1 error as a function of standard error.

**Usage**

```
plotExpectedType1Error(
  logRrNegatives,
  seLogRrNegatives,
  seLogRrPositives,
  alpha = 0.05,
  null = NULL,
  xLabel = "Relative risk",
  title,
  showCis = FALSE,
  showEffectSizes = FALSE,
  fileName = NULL
)
```

**Arguments**

<code>logRrNegatives</code>	A numeric vector of effect estimates of the negative controls on the log scale.
<code>seLogRrNegatives</code>	The standard error of the log of the effect estimates of the negative controls.
<code>seLogRrPositives</code>	The standard error of the log of the effect estimates of the positive controls.
<code>alpha</code>	The alpha (nominal type 1 error) to be used.
<code>null</code>	An object representing the fitted null distribution as created by the <code>fitNull</code> function. If not provided, a null will be fitted before plotting.
<code>xLabel</code>	If showing effect sizes, what label should be used for the effect size axis?
<code>title</code>	Optional: the main title for the plot
<code>showCis</code>	Show 95 percent credible intervals for the expected type 1 error.
<code>showEffectSizes</code>	Show the expected effect sizes alongside the expected type 1 error?
<code>fileName</code>	Name of the file where the plot should be saved, for example 'plot.png'. See the function <code>ggsave</code> in the <code>ggplot2</code> package for supported file formats.

**Details**

Creates a plot with the standard error on the x-axis and the expected type 1 error on the y-axis. The red line indicates the expected type 1 error given the estimated empirical null distribution if no calibration is performed. The dashed line indicated the nominal expected type 1 error rate, assuming the theoretical null distribution.

If standard errors are provided for non-negative estimates these will be plotted on the red line as yellow diamonds.

**Value**

A Ggplot object. Use the `ggsave` function to save to file.

**Examples**

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
positive <- sccs[sccs$groundTruth == 1, ]
plotExpectedType1Error(negatives$logRr, negatives$seLogRr, positive$seLogRr)
```

plotForest

*Create a forest plot***Description**

plotForest creates a forest plot of effect size estimates.

**Usage**

```
plotForest(
  logRr,
  seLogRr,
  names,
  xLabel = "Relative risk",
  title,
  fileName = NULL
)
```

**Arguments**

logRr	A numeric vector of effect estimates on the log scale
seLogRr	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$
names	A vector containing the names of the drugs or outcomes
xLabel	The label on the x-axis: the name of the effect estimate
title	Optional: the main title for the plot
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

**Details**

Creates a forest plot of effect size estimates (ratios). Estimates that are significantly different from 1 ( $\alpha = 0.05$ ) are marked in orange, others are marked in blue.

**Value**

A Ggplot object. Use the ggsave function to save to file.

**Examples**

```
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
plotForest(negatives$logRr, negatives$seLogRr, negatives$drugName)
```

---

plotMcmcTrace	<i>Plot the MCMC trace</i>
---------------	----------------------------

---

**Description**

Plot the MCMC trace

**Usage**

```
plotMcmcTrace(mcmcNull, fileName = NULL)
```

**Arguments**

mcmcNull	An object of type mcmcNull as generated using the fitMcmcNull function.
fileName	Name of the file where the plot should be saved, for example 'plot.png'. See the function ggsave in the ggplot2 package for supported file formats.

**Details**

Plot the trace of the MCMC for diagnostics purposes.

**Examples**

```
## Not run:
data(sccs)
negatives <- sccs[sccs$groundTruth == 0, ]
null <- fitMcmcNull(negatives$logRr, negatives$seLogRr)
plotMcmcTrace(null)

## End(Not run)
```

---

plotTrueAndObserved	<i>Plot true and observed values</i>
---------------------	--------------------------------------

---

**Description**

Plot true and observed values, for example from a simulation study.

**Usage**

```
plotTrueAndObserved(
  logRr,
  seLogRr,
  trueLogRr,
  xLabel = "Relative risk",
  title,
  fileName = NULL
)
```

## Arguments

<code>logRr</code>	A numeric vector of effect estimates on the log scale.
<code>seLogRr</code>	The standard error of the log of the effect estimates. Hint: often the standard error = $(\log(\text{lower bound 95 percent confidence interval}) - \log(\text{effect estimate})) / \text{qnorm}(0.025)$ .
<code>trueLogRr</code>	A vector of the true effect sizes.
<code>xLabel</code>	The label on the x-axis: the name of the effect estimate.
<code>title</code>	Optional: the main title for the plot
<code>fileName</code>	Name of the file where the plot should be saved, for example 'plot.png'. See the function <code>ggsave</code> in the <code>ggplot2</code> package for supported file formats.

## Details

Creates a forest plot of effect size estimates (ratios). Estimates that are significantly different from the true value ( $\alpha = 0.05$ ) are marked in orange, others are marked in blue.

## Value

A Ggplot object. Use the `ggsave` function to save to file.

## Examples

```
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotTrueAndObserved(data$logRr, data$seLogRr, data$trueLogRr)
```

---

sccs

*Incidence rate ratios from Self-Controlled Case Series*


---

## Description

Incidence rate ratios from Self-Controlled Case Series

## Usage

```
data(sccs)
```

## Format

A data frame with 46 rows and 4 variables:

- drugName** Name of the drug
- groundTruth** Whether the drug is a positive (1) or negative (0) control
- logRr** The log of the incidence rate ratio
- seLogRr** The standard error of the log of the incidence rate ratio

## Details

A dataset containing the incidence rate ratios (and standard errors) produced using a Self-Controlled Case Series (SCCS) design. The outcome is upper GI bleeding, the drug of interest (groundTruth = 1) is sertraline. Also included are 45 negative control drugs, for which we believe there to be no causal relation with upper GI bleeding. We used a database of medical records from general practices in the USA, the General Electric (GE) Centricity database, which contains data on 11.2 million subjects. We restricted on study period (start of 1990 through November 2003), age requirements (18 years or older), available time prior to event (180 days), and risk definition window (30 days following the prescription). Time 30 days prior to the first prescription was removed to account for possible contra-indications. Cases of upper GI bleeding were identified on the basis of the occurrence of ICD-9 diagnosis codes in the problem list. These codes pertain to esophageal, gastric, duodenal, peptic, and gastrojejunal ulceration, perforation, and hemorrhage, as well as gastritis and non-specific gastrointestinal hemorrhage. For more information on this set see Schuemie et al (2014).

## References

Schuemie MJ, Ryan PB, Dumouchel W, Suchard MA, Madigan D. Interpreting observational studies: why empirical calibration is needed to correct p-values. *Statistics in Medicine* 33(2):209-18,2014

---

simulateControls	<i>Simulate (negative) controls</i>
------------------	-------------------------------------

---

## Description

Simulate (negative) controls

## Usage

```
simulateControls(
  n = 50,
  mean = 0,
  sd = 0.1,
  seLogRr = runif(n, min = 0.01, max = 0.2),
  trueLogRr = 0
)
```

## Arguments

n	Number of controls to simulate.
mean	The mean of the error distribution (on the log RR scale).
sd	The standard deviation of the error distribution (on the log RR scale).
seLogRr	The standard error of the log of the relative risk. This is recycled for the controls. The default is to sample these from a uniform distribution.
trueLogRr	The true relative risk (on the log scale) used to generate these controls. This is recycled for the controls.

## Details

Generate point estimates given known true effect sizes and standard errors



**Examples**

```
data <- simulateControls(n = 50 * 3, mean = 0.25, sd = 0.25, trueLogRr = log(c(1, 2, 4)))
plotTrueAndObserved(data$logRr, data$seLogRr, data$trueLogRr)
```

---

southworthReplication *Relative risks from an unadjusted new-user cohort design*

---

**Description**

Relative risks from an unadjusted new-user cohort design

**Usage**

```
data(southworthReplication)
```

**Format**

A data frame with 174 rows and 4 variables:

**outcomeName** Name of the outcome

**trueLogRr** The log of the true effect size. Only provided for negative and positive controls, is NA for the outcome of interest (GI bleeding).

**logRr** The log of the incidence rate ratio

**seLogRr** The standard error of the log of the incidence rate ratio

**Details**

A dataset containing the incidence rate ratios (and standard errors) produced using a new-user cohort design that compares new-users of dabigatran to new-users of warfarin for the outcome of GI hemorrhage. The dataset includes estimates both for the outcome of interest as well as negative and positive control outcomes. Subjects are required to have 183 days of continuous observation prior to initiating treatment, a prior diagnosis of atrial fibrillation, and are required to have no prior exposure to either dabigatran or warfarin. The study computes an incidence rate ratio without any adjustment for confounders. Time at risk is defined as the time on the drug. The original study (Southworth 2013) uses the 'Mini-Sentinel Database'. For our replication, we use the Optum databases since both databases are US insurance claims databases. We analyzed 5,982 dabigatran-exposed and 19,155 warfarin-exposed subjects. For more information on this set see Schuemie et al (2017).

**References**

Schuemie MJ, Hripcsak G, Ryan PB, Madigan D, Suchard MA. Empirical confidence interval calibration for population-level effect estimation studies in observational healthcare data. *Proc Natl Acad Sci U S A*. 2018 Mar 13;115(11):2571-2577

Southworth MR, Reichman ME, Unger EF. Dabigatran and postmarketing reports of bleeding. *N Engl J Med* 368(14):1272-1274, 2013

# Index

## \*Topic **datasets**

- caseControl, [4](#)
- cohortMethod, [5](#)
- grahamReplication, [12](#)
- sccs, [23](#)
- southworthReplication, [25](#)

- calibrateConfidenceInterval, [2](#)
- calibrateP, [3](#)
- caseControl, [4](#)
- cohortMethod, [5](#)
- computeTraditionalCi, [6](#)
- computeTraditionalP, [6](#)
- convertNullToErrorModel, [7](#)

- EmpiricalCalibration, [8](#)
- evaluateCiCalibration, [8](#), [15](#), [18](#)

- fitMcmcNull, [7](#), [9](#)
- fitNull, [7](#), [10](#)
- fitSystematicErrorModel, [2](#), [7](#), [11](#)

- grahamReplication, [12](#)

- plotCalibration, [13](#)
- plotCalibrationEffect, [14](#)
- plotCiCalibration, [15](#)
- plotCiCalibrationEffect, [16](#)
- plotCiCoverage, [17](#)
- plotErrorModel, [18](#)
- plotExpectedType1Error, [19](#)
- plotForest, [21](#)
- plotMcmcTrace, [22](#)
- plotTrueAndObserved, [22](#)

- sccs, [23](#)
- simulateControls, [24](#)
- southworthReplication, [25](#)